
Bolin He, A53316428, Nov 19 2019

Table of Contents

Q1	1
Q2	5

Q1

```
close all;
clear all;
clc;

img = imread('geisel.jpg');
img = rgb2gray(img);
CED(img,70)
fprintf('I choose the thresholding value as 70.')
```



```
% function [ ] = CED(img,te)
% % Smoothing
% I = double(img);
% k = 1/159*[2 4 5 4 2;4 9 12 9 4;5 12 15 12 5;4 9 12 9 4; 2 4 5 4 2];
% I_filter = imfilter(I,k);
%
% % Finding Gradients
% kx = [-1 0 1;-2 0 2;-1 0 1];
% ky = [-1 -2 -1;0 0 0;1 2 1];
%
% Gx = imfilter(I_filter,kx);
% Gy = imfilter(I_filter,ky);
% G = sqrt(Gx.^2+Gy.^2);
% G_theta = atan(Gy./Gx);
%
% % Non-maximum Suppression (NMS)
% [row,col] = size(I_filter);
%
% for i = 1:row
%     for j =1:col
%         if G_theta(i,j) <= pi/8 && G_theta(i,j) >=-pi/8
%             G_theta(i,j) = 0;
%         elseif G_theta(i,j) <= 3/8*pi && G_theta(i,j) > pi/8
%             G_theta(i,j) = pi/4;
%         elseif G_theta(i,j) < -pi/8 && G_theta(i,j) >= -3/8*pi
%             G_theta(i,j) = -pi/4;
%         elseif G_theta(i,j) < -3/8*pi && G_theta(i,j) >= -1/2*pi
%             G_theta(i,j) = pi/2;
%         elseif G_theta(i,j) <= 1/2*pi && G_theta(i,j) > -3/8*pi
%             G_theta(i,j) = pi/2;
```

```


%         end
%     end
% end
%
% I_pad = padarray(G,[1,1],0,'both');
% G_theta = padarray(G_theta,[1,1],0,'both');
% [row_pad,col_pad] = size(I_pad);
% I_nms = zeros(row_pad,col_pad);
%
% for i = 2:row_pad-1
%     for j = 2:col_pad-1
%         if G_theta(i,j) == 0
%             if I_pad(i,j+1) > I_pad(i,j) || I_pad(i,j-1) >
I_pad(i,j)
%                 I_nms(i,j) = 0;
%             else
%                 I_nms(i,j) = I_pad(i,j);
%             end
%         elseif G_theta(i,j) == 1/4*pi
%             if I_pad(i+1,j-1) > I_pad(i,j) || I_pad(i-1,j+1) >
I_pad(i,j)
%                 I_pad(i,j) = 0;
%             else
%                 I_nms(i,j) = I_pad(i,j);
%             end
%         elseif G_theta(i,j) == -1/4*pi
%             if I_pad(i-1,j-1) > I_pad(i,j) || I_pad(i+1,j+1) >
I_pad(i,j)
%                 I_pad(i,j) = 0;
%             else
%                 I_nms(i,j) = I_pad(i,j);
%             end
%         elseif G_theta(i,j) == pi/2
%             if I_pad(i-1,j) > I_pad(i,j) || I_pad(i+1,j) >
I_pad(i,j)
%                 I_pad(i,j) = 0;
%             else
%                 I_nms(i,j) = I_pad(i,j);
%             end
%         end
%     end
% end
% end
%
% % Thresholding
% I_edge = zeros(row_pad,col_pad);
% for i = 1:row_pad
%     for j = 1:col_pad
%         if I_nms(i,j) < te
%             I_edge(i,j) = 0;
%         else
%             I_edge(i,j) = I_nms(i,j);
%         end
%     end
% end

```

```
% end
%
%
% G = imshow(uint8(G));
% title('The original gradient magnitude image')
% figure
% I_nms = imshow(uint8(I_nms));
% title('The image after NMS')
% figure
% I_edge = imshow(uint8(I_edge));
% title('The image after thresholding')
% end
```

I choose the thresholding value as 70.


The original gradient magnitude image



The image after NMS



The image after thresholding



Q2

```
close all;
clear all;
clc;

%%%% Image Car %%%%
car = imread('Car.tif');
imagesc(car);
colormap(gray);
colorbar;
title('Original Image');

car = double(car);
car_pad = padarray(car,[133,172],0,'both');
imFFT = fft2(car_pad);
imFFT = fftshift(imFFT);
figure
imagesc(-256:255,-256:255,log(abs(imFFT)));
colorbar;
title('Image after FFT')
xlabel('u');
ylabel('v');

% Notch Filter
D0 = 30;
n = 4;
Position = [-90,-167;-90,-82;-82,87;-82,173];
uk = Position(:,1);
vk = Position(:,2);
H_NR = 1;
[u,v] = meshgrid(-256:255);
for k = 1:4
    Dk1=sqrt(((u-uk(k)).^2)+(v-vk(k)).^2));
    Dk2=sqrt(((u+uk(k)).^2)+(v+vk(k)).^2));
    H_NR=H_NR.*(1./(1+(D0./Dk1).^(2*n))).*(1./(1+(D0./Dk2).^(2*n)));
end
figure
imagesc(-256:255,-256:255,abs(H_NR));
colorbar;
colormap(gray);
title('Butterworth Notch Reject Filter');
xlabel('u');
ylabel('v');

% Apply Filter
figure
imagesc(-256:255,-256:255,log(abs(imFFT.*H_NR)));
title('Image after FFT with Notch Filter');
colorbar;
xlabel('u');
ylabel('v');
```

```
imfilter = uint8(iff2(iff2shift(imFFT.*H_NR)));
figure
imagesc(imfilter(134:379,173:340))
colormap(gray);
title('Filtered Image')
colorbar;

%%%% Image Street %%%%
street = imread('Street.png');
figure;
imagesc(street);
colorbar;
colormap(gray);
title('Original Image');

street = double(street);
street_pad = padarray(street,[180,153],0,'post');
imFFT2 = fftshift(fft2(street_pad));
figure;
imagesc(-256:255,-256:255,log(abs(imFFT2)));
colorbar;
xlabel('u');
ylabel('v');

% Notch Filter
D0 = 20;
n = 3;
Position2 = [0,164;-168,0];
uk = Position2(:,1);
vk = Position2(:,2);
H_NR2 = 1;
[u,v] = meshgrid(-256:255);
for k = 1:2
    Dk1=sqrt(((u-uk(k)).^2)+(v-vk(k)).^2));
    Dk2=sqrt(((u+uk(k)).^2)+(v+vk(k)).^2));
    H_NR2=H_NR2.*(1./(1+(D0./Dk1).^(2*n))).*(1./(1+(D0./Dk2).^(2*n)));
end
figure
imagesc(-256:255,-256:255,abs(H_NR2));
colorbar;
colormap(gray);
title('Butterworth Notch Reject Filter');
xlabel('u');
ylabel('v');

% Apply Filter
figure
imagesc(-256:255,-256:255,log(abs(imFFT2.*H_NR2)));
title('Image after FFT with Notch Filter');
colorbar;
xlabel('u');
ylabel('v');

imfilter2 = uint8(iff2(iff2shift(imFFT2.*H_NR2)));
```

```
figure
imagesc(imfilter2(1:332,1:359))
colormap(gray);
title('Filtered Image')
colorbar;

%
Para1 = {'n';'D0';'u1';'v1';'u2';'v2';'u3';'v3';'u4';'v4'};
Para2 = {'n';'D0';'u1';'v1';'u2';'v2'};
num1 = [4;30;-90;-167;-90;-82;-82;87;-82;173];
num2 = [3;20;0;164;-168;0];
Parameter_in_2i = table(categorical(Para1),num1)
Parameter_in_2ii = table(categorical(Para2),num2)
```

Parameter_in_2i =

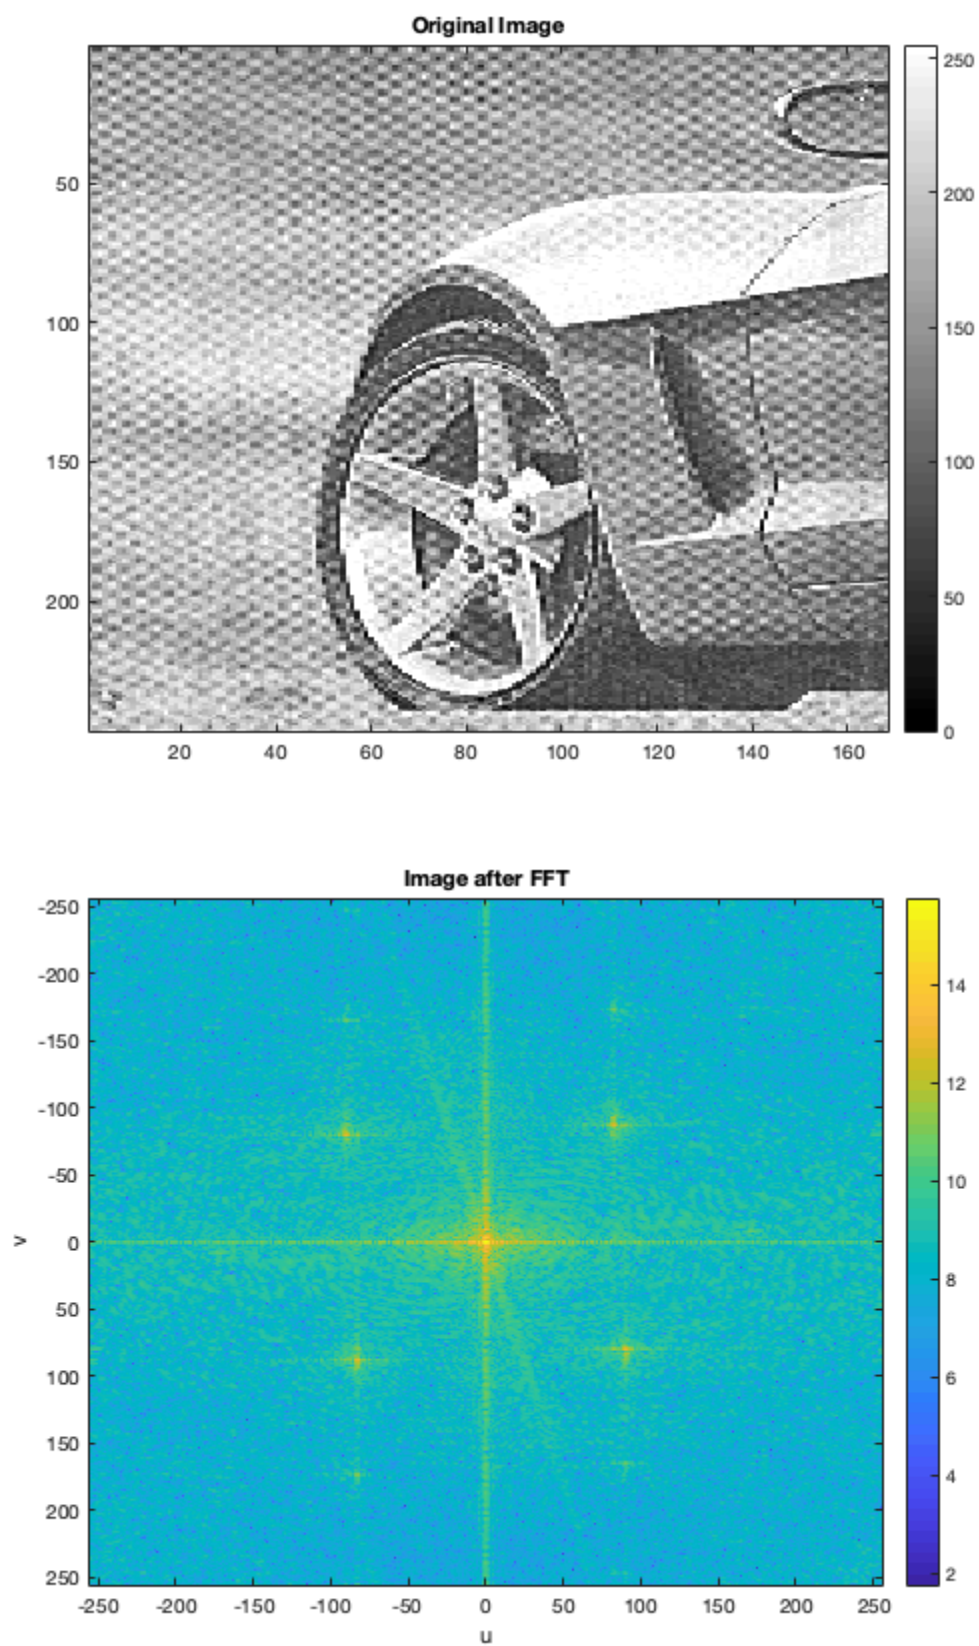
10×2 table

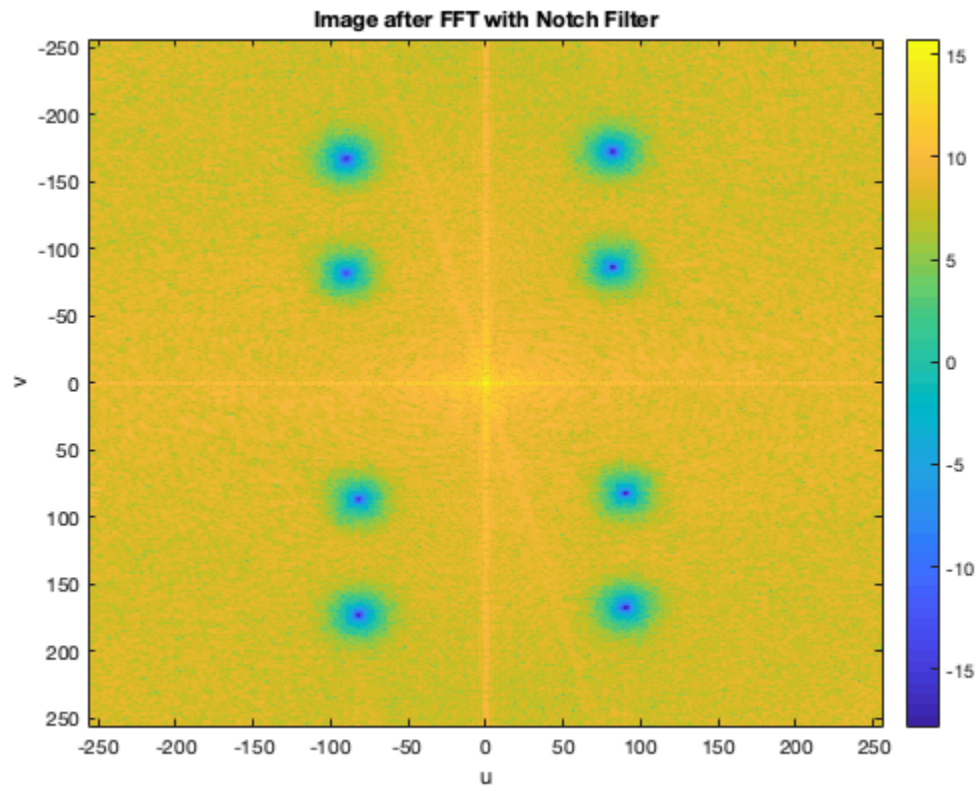
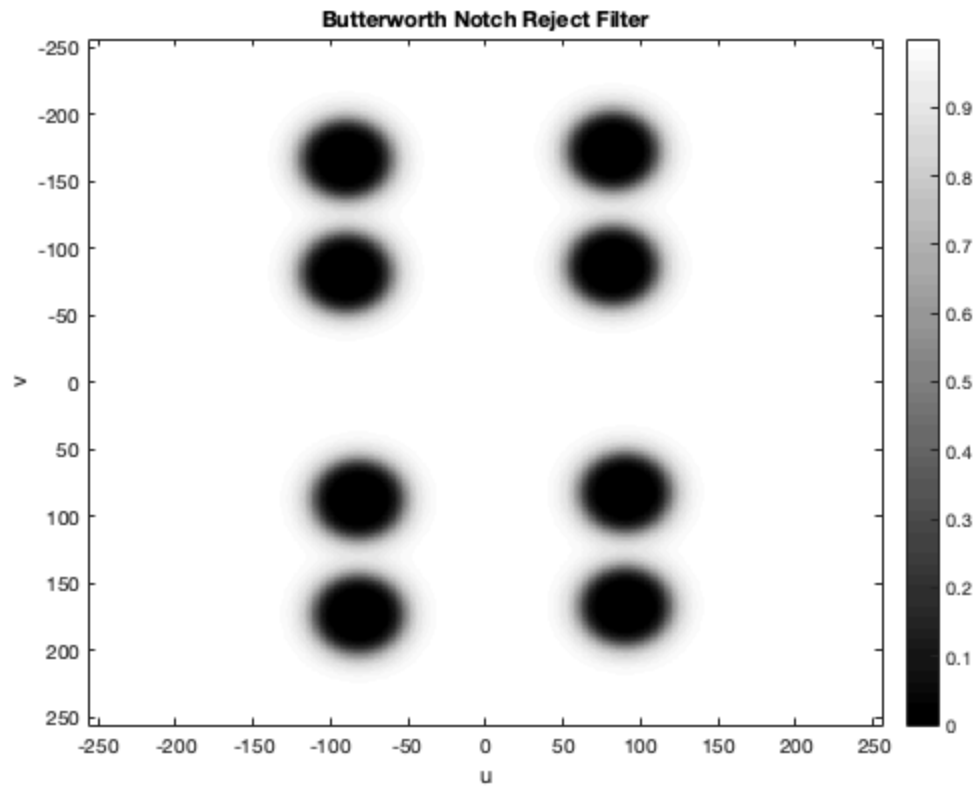
<i>Var1</i>	<i>num1</i>
<i>n</i>	4
<i>D0</i>	30
<i>u1</i>	-90
<i>v1</i>	-167
<i>u2</i>	-90
<i>v2</i>	-82
<i>u3</i>	-82
<i>v3</i>	87
<i>u4</i>	-82
<i>v4</i>	173


Parameter_in_2ii =

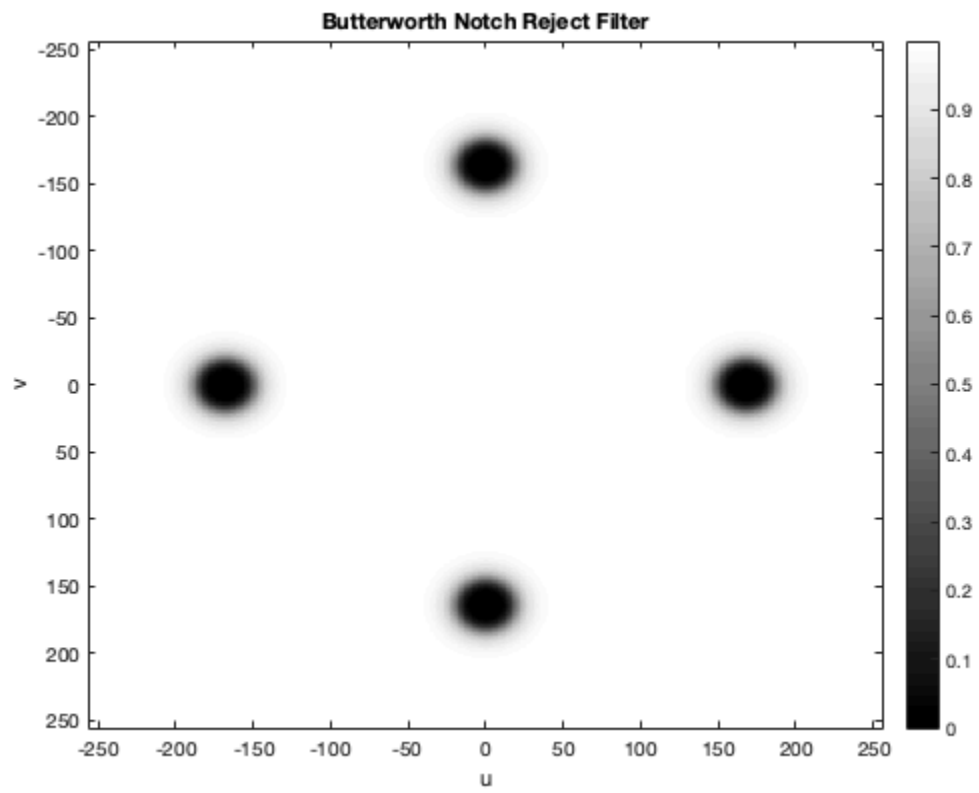

6×2 table

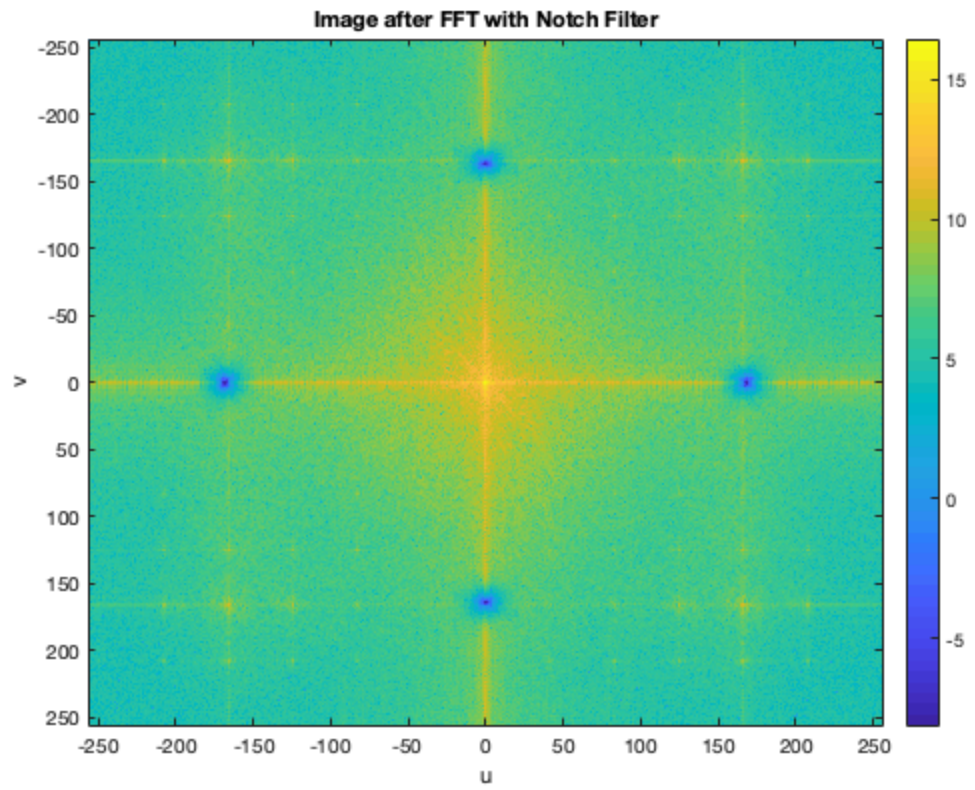
<i>Var1</i>	<i>num2</i>
<i>n</i>	3
<i>D0</i>	20
<i>u1</i>	0
<i>v1</i>	164
<i>u2</i>	-168
<i>v2</i>	0











Q3 PyTorch tutorial and questions

(ii) 50000 images are packed into 12500 batches for training.


(iii) Yes, we normalize the images. The output of torchvision datasets are PILImage images of range [0, 1]. In our example, we transform them to Tensors of normalized range [-1, 1].

(iv)

```
[1, 2000] loss: 2.209
[1, 4000] loss: 1.874
[1, 6000] loss: 1.711
[1, 8000] loss: 1.592
[1, 10000] loss: 1.525
[1, 12000] loss: 1.480
[2, 2000] loss: 1.399
[2, 4000] loss: 1.381
[2, 6000] loss: 1.351
[2, 8000] loss: 1.319
[2, 10000] loss: 1.322
[2, 12000] loss: 1.245
Finished Training
```

```
import matplotlib.pyplot as plt
plt.plot(lossf)
plt.xlabel('epoch')
plt.ylabel('loss')
```

```
Text(0, 0.5, 'loss')
```



(v)

```
class_names = trainset.classes
```

```
def visualize_model(model, num_images=4):
    was_training = model.training
    model.eval()
    images_so_far = 0
    fig = plt.figure()

    with torch.no_grad():
        for i, (inputs, labels) in enumerate(trainloader, 0):
            inputs = inputs.to(device)
            labels = labels.to(device)

            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)

            for j in range(inputs.size()[0]):
                images_so_far += 1
                ax = plt.subplot(num_images//2, 2, images_so_far)
                ax.axis('off')
                ax.set_title('predicted: {}'.format(class_names[preds[j]]))
                imshow(inputs.cpu().data[j])

            if images_so_far == num_images:
                model.train(mode=was_training)
                return
    model.train(mode=was_training)
```

```
visualize_model(net)
```


predicted: ['airplane']




predicted: ['automobile']



predicted: ['bird']



predicted: ['cat']



(vi)

```
dataiter = iter(trainloader)
input_image, label = dataiter.next()
data = net.conv1(input_image).cpu().data.numpy()
data = data[0];
```

```
kernel_num = data.shape[0]

fig, axes = plt.subplots(ncols=kernel_num, figsize=(2*kernel_num, 2))

for col in range(kernel_num):
    axes[col].imshow(data[col, :, :])
plt.show()
```

