

---

# Bolin He, PID: A53316428, Hw04

Dec 3,2019

## Q1 Hough Transform

```
clear all;
clc;
% i
% function [rho,H] = HoughTransform(img)
%     % H: HT transform matrix
%
%     % Initialization
%     [x,y] = size(img);
%     theta = [-90:90];
%     len = ceil(sqrt((x-1)^2+(y-1)^2));
%     [M,N] = find(img);
%     rho = zeros(length(x),2*90+1);
%     H = zeros(2*len+1,2*90+1);
%
%     % Calcualte rho
%     for i = 1:length(M)
%         rho(i,:) =
%             round((M(i)-1).*cosd(theta)+(N(i)-1).*sind(theta));
%     end
%
%     % Calculate H
%     for i = 1:length(M)
%         for j = 1:181
%             x_idx = rho(i,j)+len;
%             H(x_idx,j) = H(x_idx,j)+1;
%         end
%     end
%
%     Hs = H./max(H(:));
%     imshow(Hs,[],'XData',theta,'YData',-
%     len:len,'InitialMagnification','fit');
%     xlabel('\theta'), ylabel('\rho');
%     colorbar; colormap(gray);
%     axis on, axis normal
%
% end

% ii
test_img = [1 0 0 0 0 0 0 0 0 0 1
            0 0 0 0 0 0 0 0 0 0 0
            0 0 0 0 0 0 0 0 0 0 0
            0 0 0 0 0 0 0 0 0 0 0
            0 0 0 0 0 0 0 0 0 0 0
            0 0 0 0 0 1 0 0 0 0 0
            0 0 0 0 0 0 0 0 0 0 0]
```

```
0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0  
1 0 0 0 0 0 0 0 0 0 1];  
  
imshow(test_img,'InitialMagnification',5000);  
title('Original image');  
colorbar;  
figure;  
HoughTransform(test_img);  
title('HT image');  
colorbar;  
  
% select point (45,0) and (-45,7)  
figure;  
hold on;  
xx = 0:10;  
y1 = round(0/sind(45)-cosd(45)/sind(45).*xx)+10;  
y2 = round(7/sind(-45)-cosd(-45)/sind(-45).*xx)+10;  
for i = 1:length(xx)  
    test_img(xx(i)+1,y1(i)+1) = 1;  
    test_img(xx(i)+1,y2(i)+1) = 1;  
end  
imshow(test_img,'InitialMagnification',5000)  
colorbar;  
title('Original image with lines')  
  
% iii  
lane = imread('lane.png');  
lane2 = double(rgb2gray(lane));  
E = edge(lane2,'sobel');  
title('Image after Sobel Filter');  
figure  
imshow(lane)  
title("Original image");  
colorbar;  
figure  
imshow(E)  
title("Binary edge image");  
colorbar;  
figure  
[rho2, H2] = HoughTransform(E);  
title("HT")  
colorbar;  
  
% Threshold  
[A,B] = find(H2 > 0.75*max(H2(:)))  
[row,col] = size(E);  
len = ceil(sqrt((row-1)^2+(col-1)^2));  
Amap = A-len;  
Bmap = B-91;  
L = zeros(length(A),2);  
figure  
imshow(lane);
```

```
hold on;

for i=1:length(A)

    L(i,:)=[round((Amap(i))*sind(Bmap(i))+1),round((Amap(i))*cosd(Bmap(i))+1)];
    j=1:col;
    k=-(1/cotd(Bmap(i)))*(j-L(i,1))+L(i,2);
    plot(j,k,'r','LineWidth',1)
    set(gca,'ydir','reverse');
    ylim([1 row]);xlim([1 col])
end

title('Original image with lines');
colorbar;

% iv
figure()
imshow(lane);
colorbar;
hold on;
% Select theta as [-30,40]U[30,40]
drive = find(abs(Bmap)>=30 & abs(Bmap)<=40);
L2 = zeros(length(A),2);
for i=1:length(drive)

    L2(i,:)=[round((Amap(drive(i)))*sind(Bmap(drive(i))+1),round((Amap(drive(i)))*cosd(Bmap(drive(i))+1));
    j=1:col;
    k=-(1/cotd(Bmap(drive(i)))*(j-L2(i,1))+L2(i,2));
    plot(j,k,'g','LineWidth',2)
    set(gca,'ydir','reverse');
    ylim([1 row]);xlim([1 col])
end
title('Driver lane only')
close all;
fprintf('We choose theta as [-30,40]U[30,40]');

A =
363
1005
1036
1147
1172
1304
1307
1332
1337
1357
1479
1479
1482
1482
1614
```

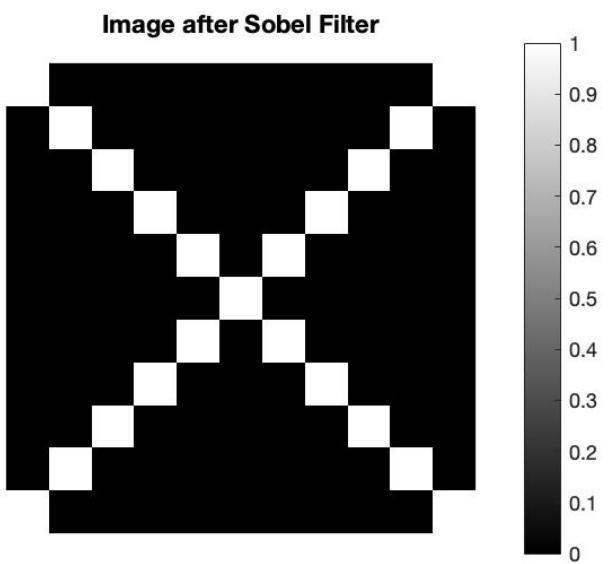
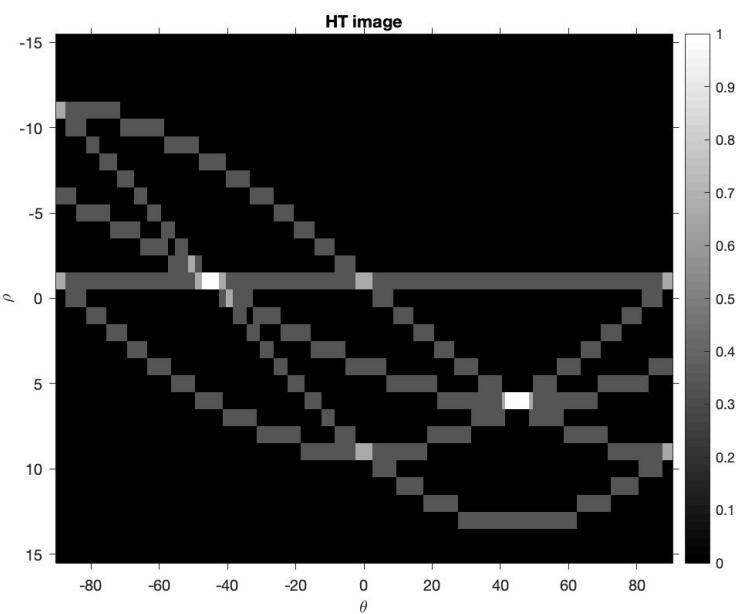
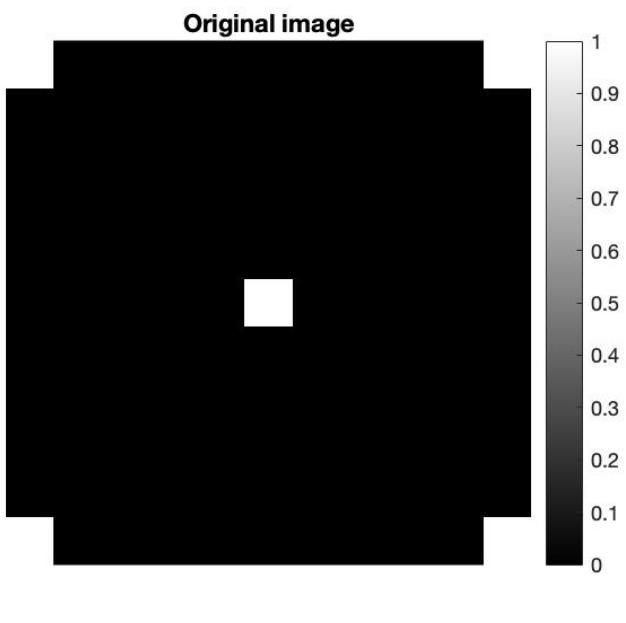
1618  
1915

$B =$

1  
53  
56  
68  
70  
84  
85  
87  
87  
91  
105  
106  
106  
107  
127  
128  
181

We choose theta as  $[-30, 40] \cup [30, 40]$

Published with MATLAB® R2018a



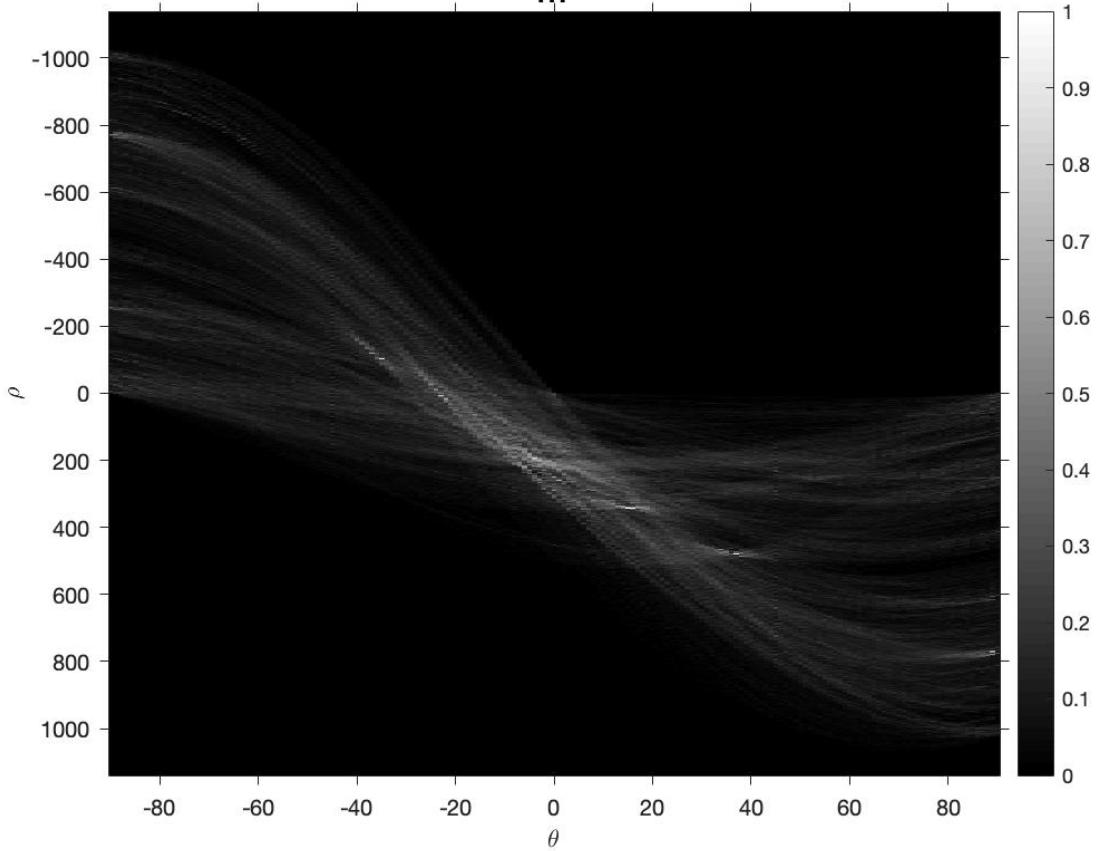
**Original image**



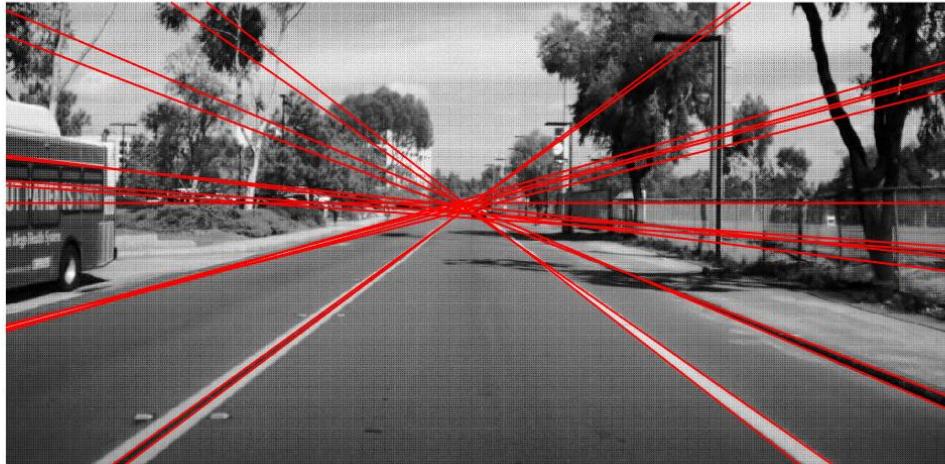
**Binary edge image**



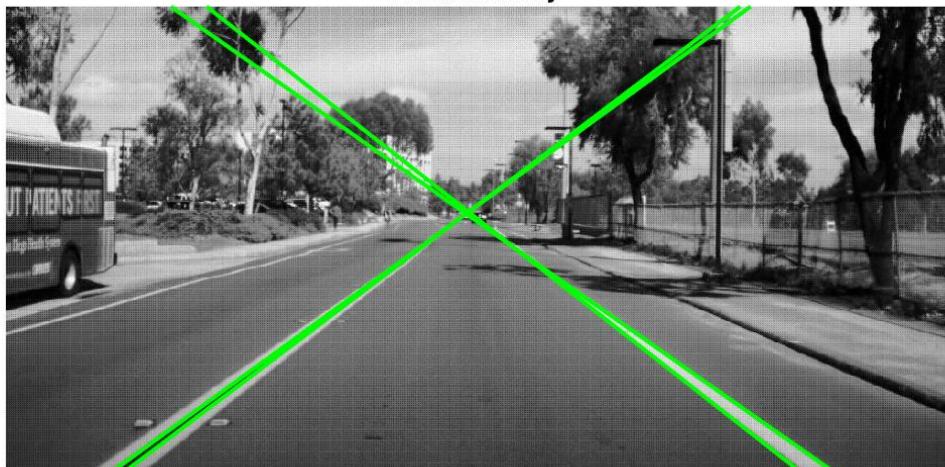
**HT**



**Original image with lines**



**Driver lane only**



---

```

% Q2 K-Means Segmentation
close all;
clear all;
clc;

im = imread('white-tower.png');
im = double(im);

features = createDataset(im);

nclusters = 7;
id = randi(size(features,1),1,nclusters);
centers = features(id,:);

[idx, centers] = KMeansCluster(features,centers);

im_seg = mapValues(im, idx);

% Print centers
fprintf('The centers of seven clusters are:\n')
table(centers)

function features = createDataset(im)
    N1 = reshape(im(:,:,1),[],1);
    N2 = reshape(im(:,:,2),[],1);
    N3 = reshape(im(:,:,3),[],1);
    features = [N1,N2,N3];
end

function [idx, centers] = kMeansCluster(features, centers)

[M,N] = size(features);
[center_row,center_col] = size(centers);

% Iterate for 100 times
    % Find the nearest center
    for i = 1:100;
        Distance = pdist2(features,centers);
        [~,idx] = min(Distance,[],2);
        centers = zeros(center_row,3);
        % Recalculate the center
        for j = 1:center_row
            center_avg = find(idx == j);
            centers(j,:) = mean(features(center_avg,:));
        end
    end
end

function [im_seg] = mapValues(im,idx)

```

---

---

```

N1 = reshape(im(:,:,1),[],1);
N2 = reshape(im(:,:,2),[],1);
N3 = reshape(im(:,:,3),[],1);
features = [N1,N2,N3];
[a,b,c] = size(im);
im_seg=zeros(a*b,c);

for i = 1:7
    idx2 = find(idx == i);
    Mean = mean(features(idx2,:));
    for j = 1:length(idx2)
        im_seg(idx2(j),:) = Mean;
    end
end

% Plot
im_seg = reshape(im_seg,[720,1280,3]);
imshow(uint8(im));
title('Original Image');
figure;
imshow(uint8(im_seg));
title('Image after segmentation');
end

```

*The centers of seven clusters are:*

```
ans =
```

```
7x1 table
```

```
centers
```

---

86.715	77.177	57.434
73.06	99.244	109.61
33.395	30.118	22.249
101.14	126.2	154.49
159.97	134.15	110.56
138.7	152.95	165.88
204.43	172.78	142.96

---

Original Image



Image after segmentation



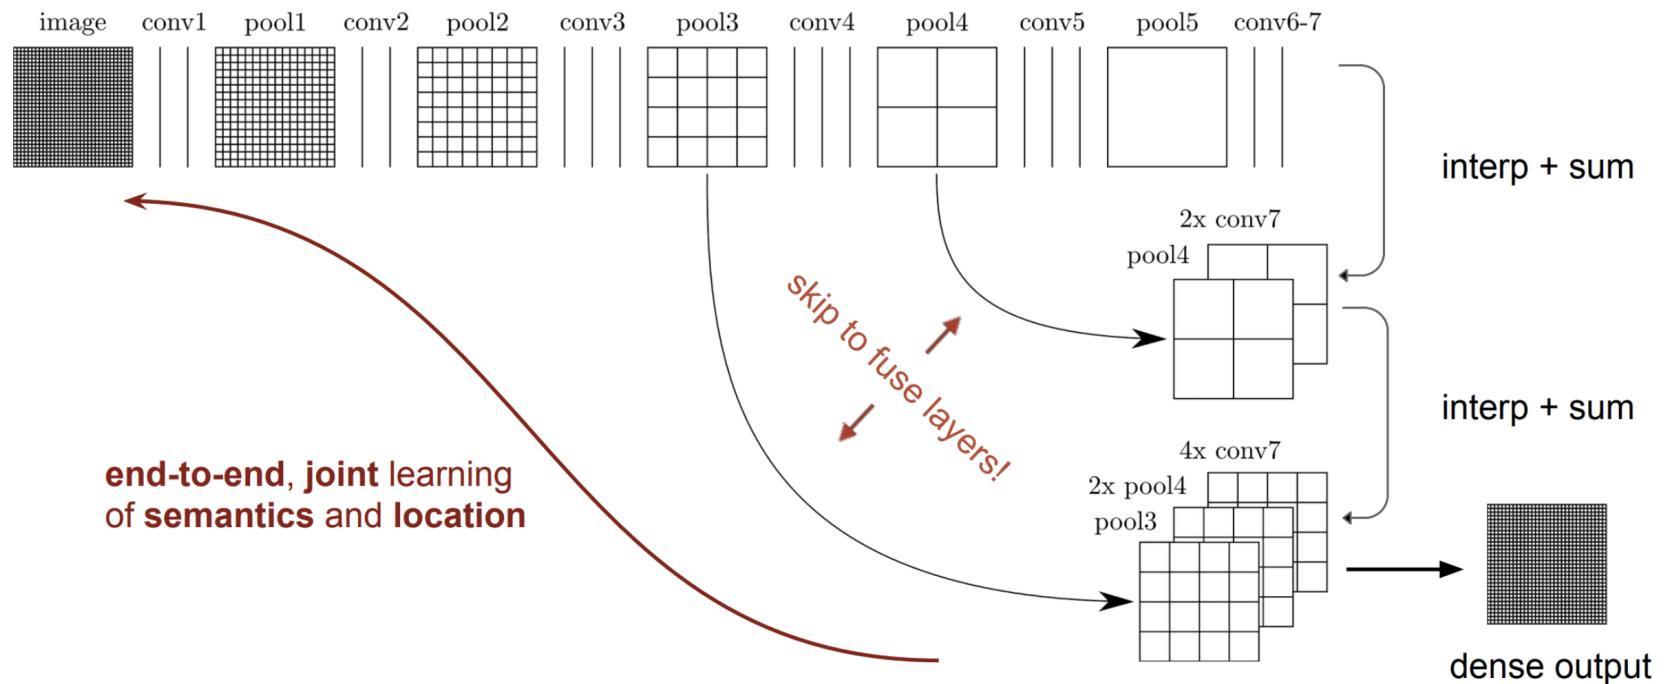
*Published with MATLAB® R2018a*

## Q1

From the figures below, we can clearly see a multi-stream network that fuses features/predictions across layers. Basically, there are variants of the FCN architecture, which mainly differ in the spatial precision of their output. For example, the figures below show the FCN-32, FCN-16 and FCN-8 variants. In the figures, convolutional layers are represented as vertical lines between pooling layers, which explicitly show the relative size of the feature maps.

As shown below, these 3 different architectures differ in the stride of the last convolution, and the skip connections used to obtain the output segmentation maps. We will use the term downsampling path to refer to the network up to conv7 layer and we will use the term upsampling path to refer to the network composed of all layers after conv7. It is worth noting that the 3 FCN architectures share the same downsampling path, but differ in their respective upsampling paths.

For the model FCN8s, it sums the 2x upsampled conv7 (with a stride 2 transposed convolution) with pool4, upsamples them with a stride 2 transposed convolution and sums them with pool3, and applies a transposed convolution layer with stride 8 on the resulting feature maps to obtain the segmentation map.



Reference: Jonathan Long, Fully Convolutional Networks for Semantic Segmentation.

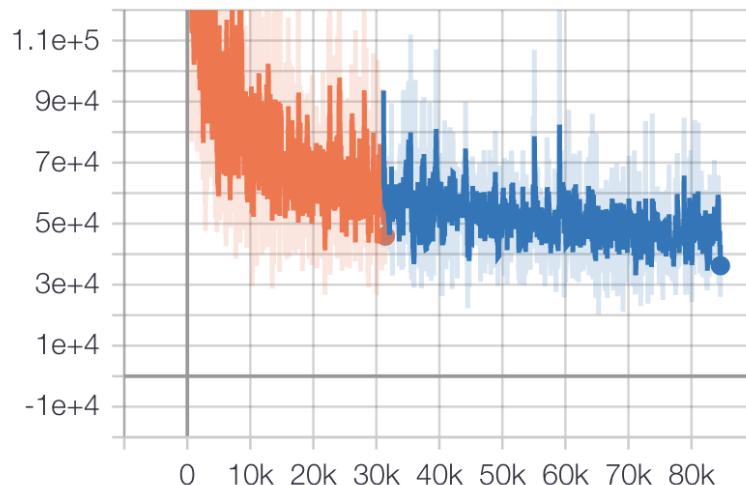
## Q2

Yes, we use the weights from Vgg16, so we did not train the model from scratch.

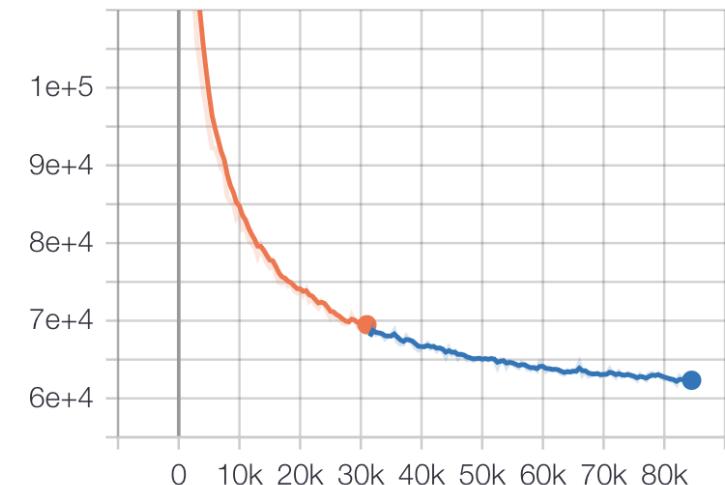
## Q3

The result from Tensorboard is:

train\_loss  
tag: loss/train\_loss



val\_loss  
tag: loss/val\_loss



## Q4

Metrics used by original paper:

- pixel accuracy
- mean accuracy
- mean IU
- frequency weighted IU

Among all the classes performance:

- Best: class 0
- Worst: class 12

## Q5

In [2]: %matplotlib notebook

```
import numpy as np
import torch
import torch.utils.data as td
import torchvision as tv
import matplotlib.pyplot as plt
import torch
import sys
sys.path.insert(0, "hw4data/ece253_pytorch_semseg_hw_design_v3/")

from ptsemseg.loader import cityscapes_loader as city
from ptsemseg.models import fcn
from ptsemseg.models import get_model
from ptsemseg.loader import get_loader
from ptsemseg.metrics import runningScore
from ptsemseg.utils import convert_state_dict
from ptsemseg.utils import recursive_glob
from PIL import Image
import yaml

torch.backends.cudnn.benchmark = True
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda

In [3]: # Setup Model

```
with open(r'hw4data/ece253_pytorch_semseg_hw_design_v3/runs/fcn8s_cityscapes/20'
          cfg = yaml.load(file, Loader=yaml.FullLoader)

state_path = 'hw4data/ece253_pytorch_semseg_hw_design_v3/runs/fcn8s_cityscapes/
```

```
In [4]: data_loader = get_loader(cfg["data"]["dataset"])
data_path = cfg["data"]["path"]

t_loader = data_loader(
    data_path,
    is_transform=True,
    split=cfg["data"]["train_split"],
    img_size=(cfg["data"]["img_rows"], cfg["data"]["img_cols"]))
)

v_loader = data_loader(
    data_path,
    is_transform=True,
    split=cfg["data"]["val_split"],
    img_size=(cfg["data"]["img_rows"], cfg["data"]["img_cols"]),
)

ts_loader = data_loader(
    data_path,
    is_transform=True,
    split="test", test_mode=True,
    img_size=(cfg["data"]["img_rows"], cfg["data"]["img_cols"]))
)
```

Found 2975 train images

Found 500 val images

Found 1525 test images

```
In [5]: trainloader = td.DataLoader(  
    t_loader,  
    batch_size=cfg[ "training" ][ "batch_size" ],  
    num_workers=cfg[ "training" ][ "n_workers" ],  
    shuffle=True  
)  
  
valloader = td.DataLoader(  
    v_loader, batch_size=cfg[ "training" ][ "batch_size" ], num_workers=cfg[ "t  
])  
  
testloader = td.DataLoader(  
    ts_loader, batch_size=cfg[ "training" ][ "batch_size" ], num_workers=cfg[ "  
])
```

```
In [6]: n_classes = t_loader.n_classes  
bs = cfg[ "training" ][ "batch_size" ]  
  
model = get_model(cfg[ 'model' ], n_classes)  
state = convert_state_dict(torch.load(state_path)[ "model_state" ])  
model.load_state_dict(state)  
model.eval()  
model.to(device);
```

Run this block multiple times to test different image

```
In [7]: # 0: visualize train set;
# 1: visualize validation set;
# 2: visualize test set
vis_mode = 1

if vis_mode == 0:
    loader, td_loader = t_loader, trainloader
elif vis_mode == 1:
    loader, td_loader = v_loader, valloader
elif vis_mode == 2:
    loader, td_loader = ts_loader, testloader

dataiter = iter(td_loader)
```

WARN: resizing labels yielded fewer classes

```
In [8]: def getImg(img_path):
    test_img = Image.open(img_path).convert('RGB')
    transform = tv.transforms.Compose([
        tv.transforms.Resize((cfg["data"]["img_rows"], cfg["data"]["img_cols"]),
        tv.transforms.ToTensor(),
        tv.transforms.Normalize((0.5,0.5,0.5), (0.5,0.5,0.5))
    ])
    return transform(test_img).unsqueeze(0)

test_img1 = getImg("hw4data/test1.jpg")
test_img2 = getImg("hw4data/test2.jpg")
test_img3 = getImg("hw4data/test3.jpg")
```

```
In [9]: # uncomment to use provided datasets
images, labels = dataiter.next()

# uncomment the following two lines to use self-taken pictures
# images = torch.cat([test_img1, test_img2, test_img3], 0)
# labels = torch.ones([images.shape[0], images.shape[2], images.shape[3]]) * 2

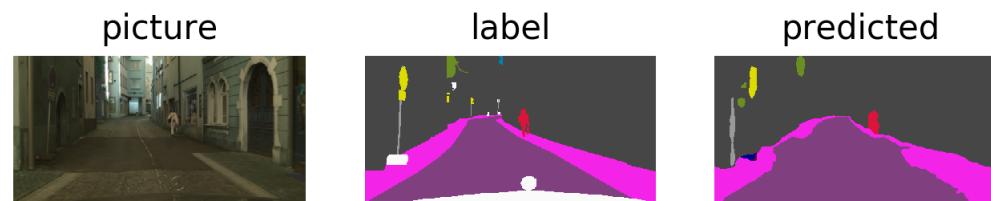
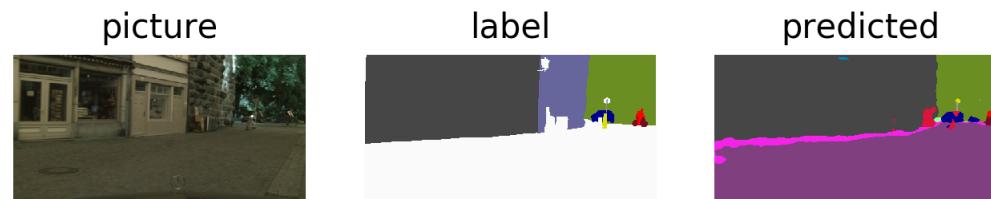
bs = images.shape[0]
imgs = images.to(device, dtype=torch.float)
pred = model(imgs).data.max(1)[1].cpu().numpy()

# show images
f, axarr = plt.subplots(bs, 3)

for j in range(bs):
    toshow = images.numpy()[j]
    toshow = np.moveaxis(toshow, [0, 1, 2], [2, 0, 1])
    axarr[j][0].imshow(toshow)
    axarr[j][1].imshow(loader.decode_segmap(labels.numpy()[j]))
    axarr[j][2].imshow(loader.decode_segmap(pred[j]))

    axarr[j][0].axis('off')
    axarr[j][1].axis('off')
    axarr[j][2].axis('off')
    axarr[j][0].set_title('picture')
    axarr[j][1].set_title('label')
    axarr[j][2].set_title('predicted')

plt.show()
```



**Q6**

```
In [10]: images = torch.cat([test_img1, test_img2, test_img3], 0)
labels = torch.ones([images.shape[0], images.shape[2], images.shape[3]]) * 250

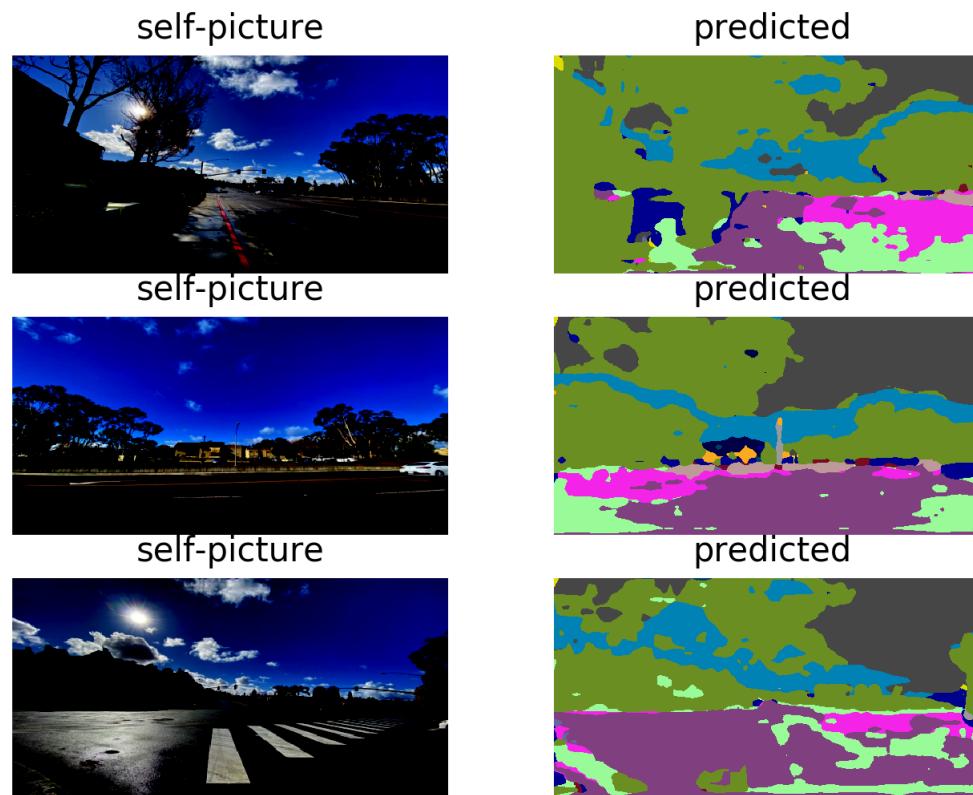
bs = images.shape[0]
imgs = images.to(device, dtype=torch.float)
pred = model(imgs).data.max(1)[1].cpu().numpy()

# show images
f, axarr = plt.subplots(bs, 2)

for j in range(bs):
    toshow = images.numpy()[j]
    toshow = np.moveaxis(toshow, [0, 1, 2], [2, 0, 1])
    axarr[j][0].imshow(toshow)
    axarr[j][1].imshow(loader.decode_segmap(pred[j]))

    axarr[j][0].axis('off')
    axarr[j][1].axis('off')
    axarr[j][0].set_title('self-picture')
    axarr[j][1].set_title('predicted')

plt.show()
```



Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

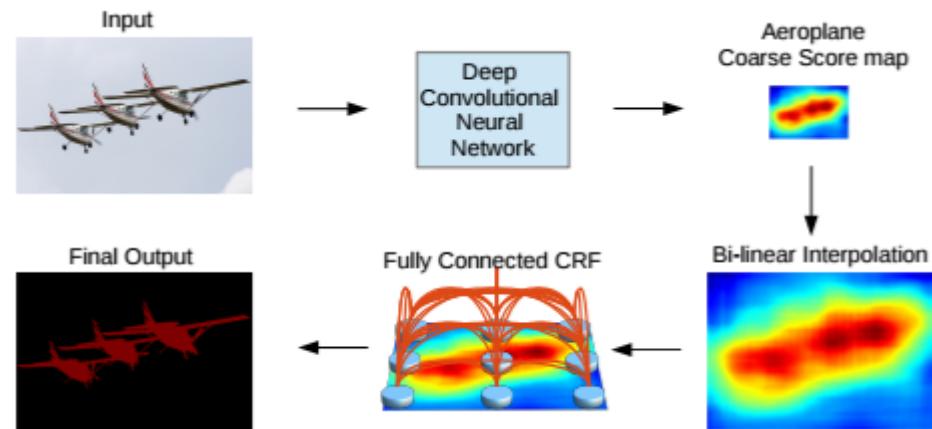
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

The result looks reasonable.

## Q7

## Improvement 1:

The upsampling process would lose some important information. We can process Fully Connected Conditional Random Fields after pixels classification, which can improve the details in figures by post processing.



Reference: Liang-Chieh Chen, Semantic image segmentation with deep convolutional nets and fully connected crfs

## Improvement 2:

We can try to achieve our goal by using RNN with CRF, which perform a better result during experiment.

Input Image      FCN-8s      DeepLab      CRF-RNN      Ground Truth



B-ground	Aero plane	Bicycle	Bird	Boat	Bottle	Bus
Car	Cat	Chair	Cow	Dining-Table	Dog	Horse
Motorbike	Person	Potted-Plant	Sheep	Sofa	Train	TV/Monitor

Reference: Shuai Zheng, Conditional random fields as recurrent neural networks

### **Improvement 3:**

We can try to use dilated convolution instead of traditional convolution in the paper. In this case, we can increase the receptive fields and retain more information for upsampling.

Reference: Fisher Yu, MULTI-SCALE CONTEXT AGGREGATION BY DILATED CONVOLUTIONS